

Dijkstra

Dijkstra's algorithm is a popular shortest-path algorithm used to find the minimum distance between a starting node and all other nodes in a weighted graph with non-negative edge weights. It works by gradually exploring the graph, always choosing the node with the smallest known distance and updating the distances of its neighboring nodes. The algorithm uses a priority queue to efficiently select the next closest node to process. As it continues, Dijkstra's algorithm builds the shortest-path tree, ensuring that once the shortest distance to a node is finalized, it never changes. It is widely used in network routing, navigation systems, and various optimization problems because of its efficiency and accuracy in computing shortest paths.

Advantages / Benefits

- Efficiently finds the shortest path from a single source to all vertices.
- Guarantees optimal paths in graphs with non-negative edge weights.
- Widely used in network routing, GPS navigation, and mapping applications.
- Can be efficiently implemented using a priority queue or min-heap for large graphs.
- Works for both dense and sparse graphs when appropriate data structures are used.
- Once a vertex is processed, its shortest distance is final, simplifying path reconstruction.

Limitations

- Cannot handle negative edge weights, unlike Bellman-Ford.
- Time complexity with a simple array is $O(V^2)$, which can be inefficient for large graphs.
- Even with a priority queue, time complexity is $O((V + E) \log V)$, which may be slow in some cases.
- Does not detect negative weight cycles, so it is unsuitable for graphs containing them.

- For dynamic graphs where edges change frequently, the algorithm may need to be rerun entirely.

Steps of Dijkstra's Algorithm

1. Initialize distances

- Set the distance of the **source node** to **0**.
- Set the distance of **all other nodes** to **infinity (∞)**.

2. Mark all nodes as unvisited

Keep a set or array to track which nodes have been processed.

3. Use a priority queue (min-heap)

This helps pick the node with the **smallest current distance** efficiently.

4. Insert the source node into the priority queue

The priority queue stores pairs like (distance, node).

5. Repeat until the priority queue becomes empty:

- Extract the node with the **smallest distance** (the current node).
- If that node is already visited, skip it.
- Mark the current node as visited.

6. Relax (update) the distances of all neighbors

For each neighbor of the current node:

- Compute a new possible distance:

$$\text{newDist} = \text{distance[current]} + \text{weight(current, neighbor)}$$
- If $\text{newDist} < \text{distance[neighbor]}$:
 - Update $\text{distance[neighbor]} = \text{newDist}$
 - Push $(\text{newDist}, \text{neighbor})$ into the priority queue.

7. Continue until all reachable nodes are processed

Once a node's shortest distance is confirmed, it will not change again.

8. End with shortest distances

The distance array now contains the minimum distance from the source to every other node.

Pseudocode:

```
Dijkstra(graph, start):
```

```
    create a distance array dist[] and initialize all values to INF
```

```
    create a visited array visited[] and initialize all to false
```

```
    dist[start] = 0
```

```
    priority_queue< pair<int, int>, vector<pair<int, int>>, greater<pair<int, int>> > pq
```

```
    pq.push({0, start})
```

```
    while pq is not empty:
```

```
        (currentDist, node) = pq.top()
```

```
        pq.pop()
```

```
        if visited[node] == true:
```

```
            continue
```

```
        visited[node] = true
```

```
        for each (neighbor, weight) in graph[node]:
```

```
if dist[node] + weight < dist[neighbor]:  
    dist[neighbor] = dist[node] + weight  
    pq.push({dist[neighbor], neighbor})
```

Time Complexity:

Time Complexity: $O((V + E) \log V)$

Where:

V = Number of Vertices/Nodes

E = Number of Edges

Dijkstra Code:

<https://github.com/Hazra32/Algorithm-Problem/blob/main/Dijkstra/Basic/basic.cpp>